

Software Design Document (SDD)

XRayVision AI — Full-Stack Medical Diagnostic Platform

Field	Value
Document	Software Design Document / System Design Specification
Version	1.0
Standard	IEEE Std 1016 (adapted)
Companion documents	SRS · Test Cases · User Manual
Institution	Minhaj University Lahore — School of Software Engineering
Programme	BSSE, 8th Semester — Final Year Project
Students	Muhammad Ali Raza (2022F-MUL-BSSWE-017), Hamza Afzal (2022F-MUL-BSSWE-027)
Supervisor	Maam Misbah — Lecturer, School of Software Engineering

Table of Contents

- 1. [Introduction](#)
- 2. [System Architecture](#)
- 3. [Data Design](#)
- 4. [Data Flow Diagrams](#)
- 5. [Component Design](#)
- 6. [Interface Design](#)
- 7. [Behavioural Design](#)
- 8. [Algorithm Design](#)
- 9. [Security Design](#)
- 10. [Error Handling and Degradation Design](#)
- 11. [Deployment Design](#)
- 12. [Design Decisions and Rationale](#)

1. Introduction

1.1 Purpose

This document describes **how** XRayVision AI is built. It translates each requirement in the [SRS](#) into concrete architecture, data structures, module boundaries, interfaces and algorithms. It is written for the evaluation committee and for any developer who must extend or maintain the system.

1.2 Scope

The design covers three deployable units — the React frontend, the FastAPI backend, and the Supabase data tier — plus the integrations with OpenRouter, Hugging Face Hub and OpenStreetMap.

1.3 Design Goals

Goal	How the design achieves it
Separation of concerns	Strict three-layer backend: routers handle HTTP only, services hold business and ML logic, utils own persistence.
Graceful degradation	Every external dependency is wrapped so a failure downgrades the result instead of destroying it.
Free-tier viability	Lazy/background model loading, CPU-only inference, rate limiting, and a single small container.
Statelessness	JWT authentication; the backend keeps no session state, so it can scale to zero and cold-start safely.
Single source of truth for I/O	One HTTP client module on the frontend, one Supabase wrapper on the backend.
Safety by construction	Closed vocabularies for urgency and severity, confidence tiering in the UI, and disclaimers rendered from the shell rather than per page.

1.4 Definitions


See [SRS §1.3](#). Additional design-level terms:

Term	Meaning in this document
Router	A FastAPI <code>APIRouter</code> module — the HTTP boundary. Contains no ML or database code.
Service	A pure-Python module holding one unit of business or ML logic, callable without HTTP.
Routed ensemble	The dispatch strategy that sends an image to the model(s) matching its <code>scan_type</code> rather than running every model on every image.
Synthesis	The LLM step that converts structured model findings into a clinical paragraph plus actions.


2. System Architecture

2.1 Architectural Style

XRayVision AI uses a **three-tier, layered architecture** with a **client-server** deployment topology. The backend internally follows a **router** → **service** → **repository** layering.



2.2 Component Diagram




2.3 Layer Responsibilities

Layer	Module pattern	Allowed to do	Forbidden from doing
Route component	<code>src/routes/*.tsx</code>	Render UI, call hooks	Call <code>fetch</code> directly
Query hook	<code>src/hooks/use-*.ts</code>	Wrap TanStack Query around <code>api.ts</code> calls	Contain business rules
HTTP client	<code>src/lib/api.ts</code>	Build requests, attach JWT, normalise errors	Render anything

Layer	Module pattern	Allowed to do	Forbidden from doing
Router	backend/app/routers/*.py	Validate input, enforce auth and rate limits, shape responses	Run ML, write SQL
Service	backend/app/services/*.py	ML inference, LLM prompting, business rules	Know about HTTP or FastAPI
Repository	backend/app/utils/supabase_client.py	All database and storage calls	Contain domain logic

3. Data Design

3.1 Entity Relationship Diagram



3.2 Data Dictionary

3.2.1 Table profiles

Column	Type	Constraints	Description
id	UUID	PK, FK → auth.users(id) ON DELETE CASCADE	Same identifier as the auth user
full_name	TEXT	NOT NULL	Display name shown in the header and PDF reports
role	TEXT	DEFAULT 'Medical Student'	Descriptive label only — not an access-control role
avatar_url	TEXT	nullable	Public URL in the avatars bucket
settings	JSONB	DEFAULT '{}'	Theme, overlay toggles, notification preferences
created_at	TIMESTAMPTZ	DEFAULT now()	Registration time
updated_at	TIMESTAMPTZ	DEFAULT now()	Last profile edit

3.2.2 Table scans

Column	Type	Constraints	Description
id	UUID	PK, DEFAULT gen_random_uuid()	Scan identifier used in the results URL
user_id	UUID	NOT NULL, FK → profiles(id) CASCADE	Owner
scan_type	TEXT	NOT NULL, CHECK IN (chest , fracture , wound)	Selected analysis route
session_label	TEXT	nullable	User-supplied patient or session tag
notes	TEXT	nullable	Clinical notes forwarded to the LLM
image_url	TEXT	NOT NULL, DEFAULT ''	Public storage URL; empty when upload failed
urgency	TEXT	CHECK IN (critical , high , medium , low , clear)	LLM-assigned urgency
findings	JSONB	NOT NULL, DEFAULT '[]'	Array of Finding objects
agent_synthesis	TEXT	nullable	The LLM clinical paragraph
agent_actions	JSONB	DEFAULT '[]'	Recommended action strings
model_results	JSONB	DEFAULT '{}'	Diagnostics: models run, model errors, processing time
created_at	TIMESTAMPTZ	DEFAULT now()	Analysis timestamp

Embedded Finding object (inside findings JSONB)

Field	Type	Description
name	string	Pathology, fracture or wound label
confidence	float	Percentage, 0–100

Field	Type	Description
severity	string	critical / high / moderate / low
model	string	Originating model name
region	string null	Anatomical region
icd_code	string null	ICD-10 code
bbox	object null	{x, y, w, h} normalised bounding box
color	string	UI hint: destructive / warning / info

3.2.3 Table chat_sessions

Column	Type	Constraints	Description
id	UUID	PK	Session identifier
user_id	UUID	NOT NULL, FK → profiles(id) CASCADE	Owner
title	TEXT	DEFAULT 'New Chat'	Session label in the sidebar
created_at	TIMESTAMPTZ	DEFAULT now()	Session start

3.2.4 Table chat_messages

Column	Type	Constraints	Description
id	UUID	PK	Message identifier
session_id	UUID	NOT NULL, FK → chat_sessions(id) CASCADE	Parent session
role	TEXT	NOT NULL, CHECK IN (user , assistant)	Message author
content	TEXT	NOT NULL	Message body
created_at	TIMESTAMPTZ	DEFAULT now()	Message time

3.3 Indexes

Index	Table	Column(s)	Purpose
idx_scans_user_id	scans	user_id	Fast history and stats lookups per user
idx_scans_created_at	scans	created_at DESC	Reverse-chronological history ordering
idx_chat_sessions_user_id	chat_sessions	user_id	Session list per user
idx_chat_messages_session_id	chat_messages	session_id	Message thread retrieval


3.4 Object Storage Design

Bucket	Visibility	Path convention	Contents
xray-images	Public	{user_id}/{scan_id}.{ext}	Uploaded medical images
avatars	Public	{user_id}.{ext}	Profile pictures

Both buckets are created idempotently at first use by `_ensure_bucket()`. Avatar URLs are returned with a `?t={timestamp}` cache-buster so a replaced avatar appears immediately.

Design note. The buckets are public because the frontend renders images directly via `<img src>` and the backend may be cold. Paths contain an opaque UUID and no patient identifier, so a URL is unguessable, but it is not access-controlled. A production deployment should switch to private buckets with short-lived signed URLs — this is recorded as a known limitation.


3.5 Row Level Security Policy Design




Writes are performed with the **service-role key** from the backend, which has already verified the JWT and resolved the owning `user_id`. Reads are additionally constrained by RLS, so a leaked anon key cannot expose another user's data.

4. Data Flow Diagrams


4.1 DFD Level 0 — Context Diagram



4.2 DFD Level 1 — Major Processes



4.3 DFD Level 2 — Process 2.0 "Analyse Image" Decomposed



5. Component Design

5.1 Backend Package Structure


backend/app/	
├─ main.py	Application factory, lifespan, CORS, rate limiter, router
├─ config.py	Pydantic Settings – every environment variable, cached with
├─ models/	
└─ schemas.py	All Pydantic request and response models
├─ routers/	HTTP boundary – no ML, no SQL
└─ auth.py	register, login, forgot-password, me, avatar, profile, set
└─ analyze.py	POST /analyze – the routed ensemble pipeline
└─ scans.py	list, get, delete, report.pdf, export.json
└─ chat.py	POST /chat, session and message listing
└─ diet.py	POST /diet
└─ stats.py	GET /stats
└─ clinics.py	GET /clinics – Overpass query and Haversine sort
├─ services/	Business and ML logic – framework-agnostic
└─ auth_service.py	JWT create/decode, FastAPI auth dependencies
└─ image_preprocess.py	Validation, DICOM decode, CLAHE, resize, normalise
└─ chest_model.py	DenseNet121 inference, 18-class mapping
└─ fracture_model.py	YOLOv8 detection, positive-box filtering
└─ fracture_classifier.py	Image-level fracture screening
└─ wound_model.py	ViT wound classification
└─ openrouter_client.py	Shared HTTP client for OpenRouter
└─ openrouter_agent.py	Clinical synthesis prompt and response parsing
└─ chatbot_service.py	Bilingual system prompts, context window, field extraction
└─ diet_service.py	Diet prompt, guardrails, plan validation, fallback plan
└─ utils/	
└─ supabase_client.py	Every database and storage operation

5.2 Frontend Structure

```
src/
├─ router.tsx           Router creation, default options
├─ routeTree.gen.ts     Generated route tree (do not edit)
├─ styles.css           Design tokens, keyframes, theme variables
├─
├─ routes/              One file per URL
│   ├─ __root.tsx       Root layout, providers, error boundary
│   ├─ index.tsx        Public landing page
│   ├─ auth.login.tsx    Login with demo-account shortcut
│   ├─ auth.register.tsx Registration
│   ├─ auth.forgot-password.tsx Password reset request
│   ├─ dashboard.tsx     Analytics
│   ├─ analyze.tsx       Three-step upload wizard + processing overlay
│   ├─ results.$scanId.tsx Tiered findings, overlays, exports
│   ├─ history.tsx       Scan list
│   ├─ chat.tsx          Bilingual chatbot with voice input
│   ├─ diet.tsx          Diet plan generator
│   ├─ clinics.tsx       GPS clinic locator
│   ├─ profile.tsx       Profile and avatar
│   └─ settings.tsx      Theme and analysis preferences
├─
├─ components/
│   ├─ app/AppShell.tsx  Sidebar, header, disclaimer badge
│   ├─ app/AuthShell.tsx Split auth layout
│   ├─ landing/          Landing sections, Hero3D, Reveal, XrayViewer
│   ├─ ui/               Shadcn/Radix primitives
│   └─ ui-x/             Project-specific composites
├─
├─ hooks/
│   ├─ use-analyze.ts    POST /analyze mutation
│   ├─ use-scans.ts       Scan list, single scan, stats queries
│   ├─ use-chat.ts        Chat mutation and session queries
│   └─ use-mobile.tsx     Responsive breakpoint hook
├─
└─ lib/
    ├─ api.ts            Single HTTP client – every call, JWT attachment, error nor
    ├─ auth-context.tsx  Global auth state, useAuth hook
    └─ types.ts          TypeScript mirrors of the Pydantic schemas
```

└ error-capture.ts	Runtime error capture
└ utils.ts	Shared helpers

5.3 Class Diagram — Domain Model (Pydantic Schemas)



5.4 Service Module Specifications

Module	Key operations	Inputs	Outputs	Failure behaviour
image_preprocess	validate_image_file , load_image_from_bytes , load_dicom_from_bytes , CLAHE + resize	raw bytes, filename, content type	numpy array	Raises <code>ValueError</code> with a user-readable reason → router maps to HTTP 422
chest_model	_get_model (cached), inference, 18-class mapping	224×224 grayscale tensor	list of findings with ICD-10 codes	Exception recorded in <code>model_errors</code> ; pipeline continues
fracture_model	YOLOv8 predict, filter positive boxes	original-size BGR array	findings with <code>bbox</code>	Same as above
fracture_classifier	Image-level fracture screening	224×224 RGB	single finding	Disabled via <code>FRACTURE_CLASSIFIER_ENABLED</code>
wound_model	ViT classification	RGB PIL image	top-1 wound finding	Same as above
openrouter_agent	synthesize_report — build clinical prompt, parse response	findings, scan type, notes	AgentSynthesis dict	Router substitutes a fallback synthesis
chatbot_service	Language-specific system prompt, 10-message context, <code>DOCTOR_TYPE</code> / <code>HOME_REMEDIES</code> extraction	message, history, language	reply + structured fields	Error surfaced to the UI

Module	Key operations	Inputs	Outputs	Failure behaviour
diet_service	Guardrail prompt, JSON parse, 7-day validation	condition, preferences, restrictions, goals	DietPlanResponse	Validated hard-coded fallback plan
auth_service	hash_password , verify_password , create_access_token , decode_token , get_current_user_id	credentials / token	JWT / user id	Raises HTTP 401
supabase_client	insert_scan , upload_image , _ensure_bucket , profile and chat operations	domain objects	rows, URLs	Caller decides whether the failure is blocking

6. Interface Design

6.1 REST API Contract

#	Method	Path	Auth	Rate limit	Request	Response
1	POST	/auth/register	—	200/min	RegisterRequest	AuthResponse
2	POST	/auth/login	—	200/min	LoginRequest	AuthResponse
3	POST	/auth/forgot-password	—	200/min	{email}	{message}
4	GET	/auth/me	JWT	200/min	—	UserProfile
5	POST	/auth/avatar	JWT	200/min	multipart/form-data	{avatar_url}
6	PATCH	/auth/profile	JWT	200/min	ProfileUpdateRequest	UserProfile
7	PATCH	/auth/settings	JWT	200/min	SettingsUpdateRequest	UserProfile
8	POST	/analyze	JWT	10/min	multipart/form-data : file , scan_type , session_label , notes	ScanResult
9	GET	/scans	JWT	200/min	optional scan_type filter	ScanListResponse
10	GET	/scans/{id}	JWT	200/min	—	ScanResult
11	DELETE	/scans/{id}	JWT	200/min	—	{message}
12	GET	/scans/{id}/report.pdf	JWT	200/min	—	application/pdf stream
13	GET	/scans/{id}/export.json	JWT	200/min	—	raw findings JSON

#	Method	Path	Auth	Rate limit	Request	Response
14	POST	/chat	JWT	20/min	ChatRequest	ChatResponse
15	GET	/chat/sessions	JWT	200/min	—	list[ChatSession]
16	GET	/chat/sessions/{id}/messages	JWT	200/min	—	list[ChatMessage]
17	POST	/diet	JWT	10/min	DietRequest	DietPlanResponse
18	GET	/stats	JWT	200/min	—	DashboardStats
19	GET	/clinics	JWT	200/min	lat , lon , radius_km (0.5–25)	ClinicSearchResponse
20	GET	/	—	200/min	—	service metadata
21	GET	/health	—	200/min	—	{"status": "ok"}

6.2 HTTP Status Code Design


Code	Meaning in this system	Example trigger
200	Success	Any successful request
400	Malformed domain input	Empty file, invalid scan_type , duplicate email
401	Authentication failure	Missing, malformed or expired JWT; bad credentials
404	Resource not found or not owned	GET /scans/{id} for another user's scan
422	Input failed quality validation	File too small, too large, or below 200×200 px
429	Rate limit exceeded	11th /analyze call within a minute
500	Unrecoverable server error	All models failed to run

6.3 Frontend–Backend Interface Rules


Rule	Rationale
Every call passes through src/lib/api.ts	One place to attach the JWT, set the base URL, and normalise errors
API_URL has its trailing slash stripped	Prevents double-slash 404s
The token is read from local storage on each request	Survives page reloads without a server session
A 401 response clears the stored token and redirects to /auth/login	Expired-token recovery without a manual logout
TanStack Query owns caching and invalidation	After a successful analysis, the scans and stats queries are invalidated so the dashboard is fresh

7. Behavioural Design


7.1 Sequence Diagram — Registration and Login




7.2 Sequence Diagram — Image Analysis (primary flow)




7.3 Sequence Diagram — Bilingual Chat




7.4 Sequence Diagram — PDF Report Export



7.5 State Diagram — Scan Lifecycle



7.6 State Diagram — Authentication Session



8. Algorithm Design

8.1 Routed Ensemble Dispatch

Rather than running all models on every image, the pipeline routes by `scan_type`. This keeps inference cheap enough for free-tier CPU hosts and keeps the report clinically focused.

```

FUNCTION run_routed_ensemble(file_bytes, scan_type, confidence_threshold):
    findings ← []
    errors ← {}
    models_run ← []

    IF scan_type = "chest":
        TRY
            result ← await to_thread(chest_model.infer, file_bytes, confidence_threshold)
            IF length(result) > MAX_CHEST_FINDINGS: // MAX_CHEST_FINDINGS
                result ← sort(result, by confidence DESC)[0 .. 6]
            findings ← result ; models_run += "DenseNet121"
        CATCH e: errors["chest"] ← message(e)

    ELSE IF scan_type = "fracture":
        TRY
            boxes ← await to_thread(fracture_model.detect, file_bytes)
            boxes ← filter(boxes WHERE label ≠ "Not_Fracture") // negative boxes prov
            findings += boxes ; models_run += "YOLOv8"
        CATCH e: errors["fracture_detector"] ← message(e)

    IF FRACTURE_CLASSIFIER_ENABLED:
        TRY
            screen ← await to_thread(fracture_classifier.classify, file_bytes)
            findings += screen ; models_run += "FractureClassifier"
        CATCH e: errors["fracture_classifier"] ← message(e)

    ELSE IF scan_type = "wound":
        TRY
            findings ← await to_thread(wound_model.classify, file_bytes)
            models_run += "ViT"
        CATCH e: errors["wound"] ← message(e)

    RETURN findings, errors, models_run

```

Complexity. One or two model forward passes per request — $O(1)$ in the number of registered models, not $O(n)$. Inference runs in a worker thread via `asyncio.to_thread`, so the event loop stays responsive.

8.2 Image Validation

```
FUNCTION validate_image_file(bytes, filename, content_type):
    size ← length(bytes)
    IF size < 100 KB: RAISE ValueError("too small – insufficient pixel data")
    IF size > 20 MB:  RAISE ValueError("too large – maximum is 20 MB")

    is_dicom ← filename ends with ".dcm"
                OR content_type ∈ {"application/dicom", "application/octet-stream"}

    IF NOT is_dicom:
        TRY PIL.open(bytes).verify()
        CATCH RAISE ValueError("not a decodable image – upload JPEG, PNG or DICOM")

        (w, h) ← PIL.open(bytes).size
        IF w < 200 OR h < 200:
            RAISE ValueError("resolution too low – minimum 200x200 px")
```

DICOM files skip the PIL checks because `pydicom` decodes them separately at preprocessing time.

8.3 Confidence Tiering (Results Page)

```
primary    ← findings WHERE confidence ≥ 65
secondary  ← findings WHERE 50 ≤ confidence < 65
borderline ← findings WHERE confidence < 50           // collapsed by default

IF ANY finding WITH confidence < 60:
    SHOW banner "radiologist review recommended"
```

This makes uncertainty visible by construction rather than leaving the user to interpret raw percentages, satisfying **NFR-S5** and **NFR-S7**.

8.4 Haversine Distance

Used to rank clinics returned by the Overpass API.

$$d = 2R * \arcsin(\sqrt{ \sin^2(dphi/2) + \cos(phi1) * \cos(phi2) * \sin^2(dlambda/2) })$$

where $R = 6371$ km, ϕ = latitude in radians, λ = longitude in radians

where $R = 6371.0$ km, ϕ is latitude in radians and λ is longitude in radians.

```
FUNCTION find_nearby_clinics(lat, lon, radius_km):
  radius_m ← radius_km × 1000
  FOR EACH mirror IN [overpass-api.de, overpass.kumi.systems, z.overpass-api.de]:
    TRY
      elements ← HTTP POST mirror WITH query(amenity|healthcare, around radius_m)
      BREAK
    CATCH timeout, HTTP error, transport error:
      CONTINUE // try the next mirror

  clinics ← []
  FOR EACH element IN elements:
    d ← haversine(lat, lon, element.lat, element.lon)
    IF d ≤ radius_km:
      clinics += ClinicResult(name, type, address, lat, lon, d, maps_url)

  RETURN sort(clinics, by distance_km ASC)
```


8.5 Diet Plan Validation

```
plan ← LLM(prompt WITH condition guardrails)
TRY
  parsed ← JSON.parse(plan)
  IF length(parsed.plan) ≠ 7: RAISE Invalid
  FOR EACH day IN parsed.plan:
    IF day lacks breakfast OR lunch OR dinner: RAISE Invalid
  RETURN parsed
CATCH Invalid, ParseError:
  RETURN HARDCODED_VALIDATED_FALLBACK_PLAN
```

A malformed plan is never shown to the user — the design prefers a known-safe plan over a partially generated one.

9. Security Design

9.1 Authentication and Authorisation Flow



Ownership is enforced **twice** — once in the handler using the token-derived `user_id` , and again by RLS in the database. A bug in one layer does not become a data leak.

9.2 Defence-in-Depth Summary

Layer	Control	Implementation
Transport	HTTPS only in production	Platform-provided TLS
Origin	CORS allowlist	Explicit origins; localhost only in dev mode
Volume	Rate limiting	SlowAPI, keyed by client IP; per-endpoint overrides
Identity	Stateless JWT	HS256, 24-hour expiry, $\geq$ 32-char secret enforced
Input	Schema validation	Pydantic on every body; explicit file quality checks
Data	Row Level Security	Policies on all four tables, <code>auth.uid()</code> scoped
Secrets	Environment variables	Pydantic Settings; nothing committed to source control
Privacy	No patient identifiers	Storage paths keyed by opaque UUIDs


9.3 Known Security Limitations

Recorded honestly for the evaluation committee:

#	Limitation	Impact	Suggested remedy
1	Storage buckets are public	Anyone holding an image URL can view it	Private buckets plus short-lived signed URLs
2	Rate limiting is per-IP, in-memory	Resets on restart; shared NAT users share a bucket	Redis-backed limiter keyed by user id
3	No refresh-token rotation	The user must sign in again every 24 hours	Add refresh tokens
4	No audit log	Deletions are not traceable	Append-only audit table
5	No malware scanning of uploads	A crafted file is passed to image decoders	Scan uploads before decoding

10. Error Handling and Degradation Design

The central design principle: **a partial result is more useful than an error page.**




Failure	Classification	System behaviour
Invalid input	Fail fast	400/422 with a specific reason
Every model failed	Hard failure	500 with the underlying message
One of several models failed	Degrade	Continue; record in <code>model_results.model_errors</code>
LLM unavailable	Degrade	Fallback synthesis, <code>urgency = medium</code>
Storage unavailable	Degrade	Empty <code>image_url</code> , warning logged

Failure	Classification	System behaviour
Diet plan malformed	Degrade	Validated hard-coded fallback plan
All Overpass mirrors down	Soft failure	Error message; user may retry
Rate limit exceeded	Reject	429 "Rate limit exceeded. Please slow down."


11. Deployment Design

11.1 Deployment Diagram



Deployment note. The repository also contains `render.yaml1` (Render as an alternative backend host) and `wrangler.jsonc` (Cloudflare Workers as an alternative frontend host). The **live deployment of record** is Vercel for the frontend and Hugging Face Spaces for the backend; the other configurations are retained as portability evidence for constraint **C-1**.

11.2 Startup Sequence



Each model load is individually wrapped in try/except, so one unavailable model never blocks startup.

11.3 Environment Matrix

Environment	Frontend URL	Backend URL	DISABLE_PRELOAD	CORS origins
Local development	<code>http://localhost:5173</code>	<code>http://localhost:8000</code>	<code>false</code>	localhost set auto-added
Production	Vercel domain	Hugging Face Space	<code>false</code> on ≥ 2 GB hosts, <code>true</code> on 512 MB tiers	<code>FRONTEND_URL + ALLOWED_ORIGINS</code>

12. Design Decisions and Rationale

#	Decision	Alternatives considered	Rationale
D-1	Routed ensemble instead of running every model	Run all models and merge	A chest model on a wound photo produces noise. Routing cuts CPU cost roughly threefold and keeps reports clinically focused.
D-2	LLM synthesis on top of model outputs	Show raw model scores only	Raw multi-label probabilities are unreadable to a student. The LLM converts them into an urgency, a paragraph and next steps.
D-3	Fallback synthesis when the LLM fails	Return 503	Model findings are the valuable part; losing the prose should not lose the analysis.
D-4	Non-blocking image upload	Fail the request on upload error	The report is worth more than the thumbnail. <code>image_url</code> is nullable by design.
D-5	Ignore <code>Not_Fracture</code> YOLO boxes	Report them as "normal"	A detector's negative box covers a region, not the whole study. Reporting it as normal would be a clinically unsafe claim (NFR-S4).
D-6	Background model preloading	Load models during startup	Free-tier health checks time out during a synchronous multi-model load.
D-7	Stateless JWT	Server-side sessions	The backend may scale to zero and cold-start; sessions in memory would be lost.
D-8	Ownership checked in handler and RLS	Handler check only	Defence in depth — a single missed <code>WHERE</code> clause cannot leak data.
D-9	Single <code>api.ts</code> HTTP client	<code>fetch</code> inside components	One place for JWT attachment, base-URL normalisation and 401 handling.
D-10	Confidence tiering in the UI	Flat, sorted list	Uncertainty must be structurally visible, not left to the reader.
D-11	Top-6 cap on chest findings	Report all 18 probabilities	18 near-threshold labels bury the signal; 6 keeps the report readable.
D-12	Hard-coded fallback diet plan	Return the malformed LLM output	Nutrition advice for a medical condition must never be partially generated.
D-13	Three Overpass mirrors	A single endpoint	The public Overpass instances rate-limit aggressively; mirrors make the feature usable.
D-14	JSONB for <code>findings</code>	Normalised <code>findings</code> table	Findings are always read whole with their scan and never queried across scans. JSONB avoids a join with no loss of function.
D-15	Public storage buckets	Signed URLs	Simplifies rendering while the backend may be cold. Recorded as a known limitation in §9.3.

Appendix A — Design Traceability

SRS requirement	SDD section
FR-011 – FR-025a	§7.2 Sequence, §8.1 Routed ensemble, §8.2 Validation, §10 Degradation
FR-026 – FR-033	§3.2 Data dictionary, §6.1 API contract, §7.4 PDF export
FR-028a – FR-028d	§8.3 Confidence tiering
FR-038 – FR-043a	§7.3 Chat sequence, §5.4 chatbot_service
FR-044 – FR-048	§8.5 Diet validation
FR-049 – FR-052a	§8.4 Haversine and mirror fallback
FR-057 – FR-060	§11.2 Startup sequence
NFR-SEC1 – SEC9	§9 Security design
NFR-Q1 – Q2	§10 Error handling and degradation
NFR-P2, NFR-P5	§11.2 Startup, §8.1 thread-pool inference

Appendix B — Revision History

Version	Date	Author	Description
1.0	August 2026	Muhammad Ali Raza, Hamza Afzal	Initial complete SDD